

Home > Cheat sheets > SQL

SQL Window Functions Cheat Sheet

With this SQL Window Functions cheat sheet, you'll have a handy reference guide to the various types of window functions in SQL.

Oct 23, 2022 · 10 min read



Richie Cotton

Webinar & podcast host, course and book author, spends all day chit-chatting about data

TOPICS

SQL

SQL, also known as Structured Query Language, is a powerful tool to search through large amounts of data and return specific information for analysis. Learning SQL is crucial for anyone aspiring to be a data analyst, **data engineer**, or data scientist, and helpful in many other fields such as web development or marketing.

In this cheat sheet, you'll find a handy collection of information about window functions in SQL including what they are and how to construct them.

 To easily run all the example code in this tutorial yourself, you can [create a DataLab workbook](#) for free that has the dataset, SQL, and the code samples pre-configured. Great for experimenting!



SQL for Data Science

SQL Window Functions

Learn SQL online at www.DataCamp.com

> Example dataset

We will use a dataset on the sales of bicycles as a sample. This dataset includes:

The [product] table

The product table contains the types of bicycles sold, their model year, and list price.

| product_id | product_name | model_year | list_price |
|------------|------------------------------------|------------|------------|
| 1 | Trek 820 - 2016 | 2016 | 379.99 |
| 2 | Ritchey Timberwolf Frameset - 2016 | 2016 | 749.99 |
| 3 | SURLy Wednesday Frameset - 2016 | 2016 | 999.99 |
| 4 | Trek Fuel EX 8 29 - 2016 | 2016 | 2899.99 |
| 5 | Heller Shaganaw Frame - 2016 | 2016 | 1320.99 |

The [order] table

The order table contains the order_id and its date.

| order_id | order_date |
|----------|--------------------------|
| 1 | 2016-01-17T00:00:00.000Z |
| 2 | 2016-01-01T00:00:00.000Z |
| 3 | 2016-01-02T00:00:00.000Z |
| 4 | 2016-01-03T00:00:00.000Z |
| 5 | 2016-01-03T00:00:00.000Z |

The [order_items] table

The order_items table lists the orders of a bicycle store. For each order_id, there are several products sold (product_id). Each product_id has a discount value.

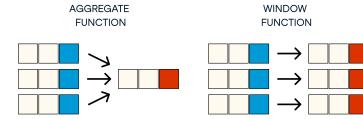
| order_id | product_id | discount |
|----------|------------|----------|
| 1 | 20 | 0.2 |
| 1 | 8 | 0.07 |
| 1 | 10 | 0.05 |
| 1 | 16 | 0.05 |
| 1 | 4 | 0.2 |
| 2 | 20 | 0.07 |

What are Window Functions?

A window function makes a calculation across multiple rows that are related to the current row. For example, a window function allows you to calculate:

- Running totals (i.e. sum values from all the rows before the current row)
- 7-day moving averages (i.e. average values from 7 rows before the current row)
- Rankings

Similar to an aggregate function (GROUP BY), a window function performs the operation across multiple rows. Unlike an aggregate function, a window function does not group rows into one single row.



> Syntax

Windows can be defined in the SELECT section of the query.

```
SELECT
    window_function() OVER(
        PARTITION BY partition_expression
        ORDER BY order_expression
        window_frame_extent
    ) AS window_column_alias
FROM table_name
```

To reuse the same window with several window functions, define a named window using the WINDOW keyword. This appears in the query after the HAVING section and before the ORDER BY section.

```
SELECT
    window_function() OVER(window_name)
    FROM table_name
    [window_name AS (
        PARTITION BY partition_expression
        ORDER BY order_expression
        window_frame_extent
    ) [ORDER BY ...]]
```

> Order by

ORDER BY is a subclause within the OVER clause. ORDER BY changes the basis on which the function assigns numbers to rows.

It is a must-have for window functions that assign sequences to rows, including RANK and ROW_NUMBER. For example, if we ORDER BY the expression 'price' on an ascending order, then the lowest-priced item will have the lowest rank.

Let's compare the following two queries which differ only in the ORDER BY clause.

```
/* Rank price from HIGH->LOW */
SELECT
    product_name,
    list_price,
    RANK() OVER(
        ORDER BY list_price DESC) rank
    FROM products
    /* Rank price from LOW->HIGH */
SELECT
    product_name,
    list_price,
    RANK() OVER(
        ORDER BY list_price ASC) rank
    FROM products
```

> Partition by

We can use PARTITION BY together with OVER to specify the column over which the aggregation is performed.

Comparing PARTITION BY with GROUP BY, we find the following similarity and difference:

- Just like GROUP BY, the OVER subclause splits the rows into as many partitions as there are unique values in a column.
- However, while the result of a GROUP BY aggregates all rows, the result of a window function using PARTITION BY aggregates each partition independently. Without the PARTITION BY clause, the result set is one single partition.

For example, using GROUP BY, we can calculate the average price of bicycles per model year using the following query:

```
SELECT
    model_year,
    AVG(list_price) avg_price
FROM products
GROUP BY model_year
```

| model_year | avg_price |
|------------|-----------|
| 2016 | 2992.21 |
| 2017 | 1277.92 |
| 2018 | 1499.41 |
| 2019 | 2689.32 |

What if we want to compare each product's price with the average price of that year? To do that, we use the AVG() window function and PARTITION BY the model year, as such.

```
SELECT
    model_year,
    product_name,
    list_price,
    AVG(list_price) OVER(
        PARTITION BY model_year
    ) avg_price
FROM products
```

Notice how the avg_price of 2018 is exactly the same whether we use the PARTITION BY clause or the GROUP BY clause.

> Window frame extent

A window frame is the selected set of rows in the partition over which aggregation will occur. Put simply, they are a set of rows that are somehow related to the current row.

A window frame is defined by a lower bound and an upper bound relative to the current row. The lowest possible bound is the first row, which is known as UNBOUNDED PRECEDING. The highest possible bound is the last row, which is known as UNBOUNDED FOLLOWING. For example, if we only want to get 5 rows before the current row, then we will specify the range using 5 PRECEDING.



> Accompanying Material

You can use this <https://bit.ly/3scZtOK> to run any of the queries explained in this cheat sheet.

Have this cheat sheet at your fingertips

 Download PDF

Associate Data Engineer in SQL

Gain practical knowledge in ETL, SQL, and data warehousing for data engineering.

Explore Track

Example dataset

We will use a dataset on the sales of bicycles as a sample. This dataset includes:

The [product] table

The product table contains the types of bicycles sold, their model year, and list price.

The [order] table

The order table contains the `order_id` and its date.

The [order_items] table

The `order_items` table lists the orders of a bicycle store. For each `order_id`, there are several products sold (`product_id`). Each `product_id` has a discount value.

What are Window Functions?

A window function makes a calculation across multiple rows that are related to the current row. For example, a window function allows you to calculate:

- Running totals (i.e. sum values from all the rows before the current row)
- 7-day moving averages (i.e. average values from 7 rows before the current row)
- Rankings

Similar to an aggregate function (`GROUP BY`), a window function performs the operation across multiple rows. Unlike an aggregate function, a window function does not group rows into one single row.

Syntax

Windows can be defined in the `SELECT` section of the query.

```
SELECT
    window_function() OVER(
        PARTITION BY partition_expression
        ORDER BY order_expression
        window_frame_extent
    ) AS window_column_alias
FROM table_name
```



POWERED BY datacamp

To reuse the same window with several window functions, define a named window using the `WINDOW` keyword. This appears in the query after the `HAVING` section and before the `ORDER BY` section.

```

SELECT
    window_function() OVER(window_name)
    FROM table_name
    [HAVING ...]
    WINDOW window_name AS (
        PARTITION BY partition_expression
        ORDER BY order_expression
        window_frame_extent
    )
    [ORDER BY ...]

```



POWERED BY datacamp

Order by

ORDER BY is a subclause within the OVER clause. ORDER BY changes the basis on which the function assigns numbers to rows.

It is a must-have for window functions that assign sequences to rows, including RANK and ROW_NUMBER. For example, if we ORDER BY the expression price on an ascending order, then the lowest-priced item will have the lowest rank.

Let's compare the following two queries which differ only in the ORDER BY clause

```

/* Rank price from LOW->HIGH */
SELECT
    product_name,
    list_price,
    RANK() OVER
        (ORDER BY list_price ASC) rank
FROM products

```



```

/* Rank price from HIGH->LOW */
SELECT
    product_name,
    list_price,
    RANK() OVER
        (ORDER BY list_price DESC) rank
FROM products

```



POWERED BY datacamp

Partition by

We can use `PARTITION BY` together with `OVER` to specify the column over which the aggregation is performed.

Comparing `PARTITION BY` with `GROUP BY`, we find the following similarity and difference:

Just like `GROUP BY`, the `OVER` subclause splits the rows into as many partitions as there are unique values in a column.

However, while the result of a `GROUP BY` aggregates all rows, the result of a window function using `PARTITION BY` aggregates each partition independently. Without the `PARTITION BY` clause, the result set is one single partition.

For example, using `GROUP BY`, we can calculate the average price of bicycles per model year using the following query.

```
SELECT
    model_year,
    AVG(list_price) avg_price
FROM products
GROUP BY model_year
```



POWERED BY databricks

What if we want to compare each product's price with the average price of that year? To do that, we use the `AVG()` window function and `PARTITION BY` the model year, as such.

```
SELECT
    model_year,
    product_name,
    list_price,
    AVG(list_price) OVER
        (PARTITION BY model_year)      avg_price
FROM products
```



POWERED BY databricks

Notice how the `avg_price` of 2018 is exactly the same whether we use the `PARTITION BY` clause or the `GROUP BY` clause.

Window frame extent

A window frame is the selected set of rows in the partition over which aggregation will occur. Put simply, they are a set of rows that are somehow related to the current row.

A window frame is defined by a lower bound and an upper bound relative to the current row. The lowest possible bound is the first row, which is known as UNBOUNDED PRECEDING . The highest possible bound is the last row, which is known as UNBOUNDED FOLLOWING . For example, if we only want to get 5 rows before the current row, then we will specify the range using 5 PRECEDING .

Accompanying Material

You can use this <https://bit.ly/3scZtOK> to run any of the queries explained in this cheat sheet.

Ranking window functions

There are several window functions for assigning rankings to rows. Each of these functions requires an ORDER BY sub-clause within the OVER clause.

The following are the ranking window functions and their description:

| Function Syntax | Function Description | Additional notes |
|------------------|--|---|
| ROW_NUMBER() | Assigns a sequential integer to each row within the partition of a result set. | Row numbers are not repeated within each partition. |
| RANK() | Assigns a rank number to each row in a partition. | <ul style="list-style-type: none"> Tied values are given the same rank. The next rankings are skipped. |
| PERCENT_RANK() | Assigns the rank number of each row in a partition as a percentage. | <ul style="list-style-type: none"> Tied values are given the same rank. Computed as the fraction of rows less than the current row, i.e., the rank of row divided by the largest rank in the partition. |
| NTILE(n_buckets) | Distributes the rows of a partition into a | <ul style="list-style-type: none"> For example, if we perform the window function NTILE(5) on a |

specified number of buckets.

table with 100 rows, they will be in bucket 1, rows 21 to 40 in bucket 2, rows 41 to 60 in bucket 3, et cetera.

CUME_DIST()

The cumulative distribution: the percentage of rows less than or equal to the current row.

- It returns a value larger than 0 and at most 1.
- Tied values are given the same cumulative distribution value.

We can use these functions to rank the product according to their prices.

```
/* Rank all products by price */
SELECT
    product_name,
    list_price,
    ROW_NUMBER() OVER (ORDER BY list_price) AS row_num,
    DENSE_RANK() OVER (ORDER BY list_price) AS dense_rank,
    RANK() OVER (ORDER BY list_price) AS rank,
    PERCENT_RANK() OVER (ORDER BY list_price) AS pct_rank,
    NTILE(75) OVER (ORDER BY list_price) AS ntile,
    CUME_DIST() OVER (ORDER BY list_price) AS cume_dist
FROM products
```



POWERED BY datacamp

Value window functions

`FIRST_VALUE()` and `LAST_VALUE()` retrieve the first and last value respectively from an ordered list of rows, where the order is defined by `ORDER BY`.

| Value window function | Function |
|---|---|
| <code>FIRST_VALUE(value_to_return) OVER (ORDER BY value_to_order_by)</code> | Returns the first value in an ordered set of values |

| | |
|---|--|
| LAST_VALUE(value_to_return) OVER (ORDER BY value_to_order_by) | Returns the last value in an ordered set of values |
| NTH_VALUE(value_to_return, n) OVER (ORDER BY value_to_order_by) | Returns the nth value in an ordered set of values. |

To compare the price of a particular bicycle model with the cheapest (or most expensive) alternative, we can use the FIRST_VALUE (or LAST_VALUE).

```
/* Find the difference in price from the cheapest alternative */
SELECT
    product_name,
    list_price,
    FIRST_VALUE(list_price) OVER (
        ORDER BY list_price
        ROWS BETWEEN
            UNBOUNDED PRECEDING
            AND
            UNBOUNDED FOLLOWING
    ) AS cheapest_price,
FROM products
```

POWERED BY  databricks

```
/* Find the difference in price from the priciest alternative */
SELECT
    product_name,
    list_price,
    LAST_VALUE(list_price) OVER (
        ORDER BY list_price
        ROWS BETWEEN
            UNBOUNDED PRECEDING
            AND
            UNBOUNDED FOLLOWING
    ) AS highest_price
FROM products
```

POWERED BY  databricks

Aggregate window functions

Aggregate functions available for GROUP BY , such as COUNT(), MIN(), MAX(), SUM(), and AVG() are also available as window functions.

| Function Syntax | Function Description |
|--|--|
| COUNT(expression) OVER (PARTITION BY partition_column) | Count the number of rows that have a non-null expression in the partition. |
| MIN(expression) OVER (PARTITION BY partition_column) | Find the minimum of the expression in the partition. |
| MAX(expression) OVER (PARTITION BY partition_column) | Find the maximum of the expression in the partition. |
| AVG(expression) OVER (PARTITION BY partition_column) | Find the mean (average) of the expression in the partition. |

Suppose we want to find the average, maximum and minimum discount for each product, we can achieve it as such.

```
SELECT
    order_id,
    product_id,
    discount,
    AVG(discount) OVER (PARTITION BY product_id) AS avg_discount,
    MIN(discount) OVER (PARTITION BY product_id) AS min_discount,
    MAX(discount) OVER (PARTITION BY product_id) AS max_discount
FROM order_items
```

POWERED BY  datacamp

LEAD, LAG

The LEAD and LAG locate a row relative to the current row.

| Function Syntax | Function Description |
|-----------------|----------------------|
| | |

| | |
|--|--|
| <code>LEAD(expression [,offset[,default_value]]) OVER(ORDER BY columns)</code> | Accesses the value stored in a row after the current row. |
| <code>LAG(expression [,offset[,default_value]]) OVER(ORDER BY columns)</code> | Accesses the value stored in a row before the current row. |

Both LEAD and LAG take three arguments:

- Expression : the name of the column from which the value is retrieved
- Offset : the number of rows to skip. Defaults to 1.
- Default_value : the value to be returned if the value retrieved is null. Defaults to NULL .

With LAG and LEAD , you must specify ORDER BY in the OVER clause.

LEAD and LAG are most commonly used to find the value of a previous row or the next row. For example, they are useful for calculating the year-on-year increase of business metrics like revenue.

Here is an example of using lag to compare this year's sales to last year's.

```
/* Find the number of orders in a year */  
WITH yearly_orders AS (  
    SELECT  
        year(order_date) AS year,  
        COUNT(DISTINCT order_id) AS num_orders  
    FROM sales.orders  
    GROUP BY year(order_date)  
)  
  
/* Compare this year's sales to last year's */  
SELECT  
    *,  
    LAG(num_orders) OVER (ORDER BY year) last_year_order,  
    LAG(num_orders) OVER (ORDER BY year) - num_orders diff_from_last_year  
FROM yearly_orders
```

POWERED BY  datacamp

Similarly, we can make a comparison of each year's order with the next year's.



```
/* Find the number of orders in a year */
WITH yearly_orders AS (
    SELECT
        year(order_date) AS year,
        COUNT(DISTINCT order_id) AS num_orders
    FROM sales.orders
    GROUP BY year(order_date)
)
```



EN

```
LEAD(num_orders) OVER (ORDER BY year) next_year_order,
LEAD(num_orders) OVER (ORDER BY year) - num_orders diff_from_next_year
FROM yearly_orders
```

POWERED BY datacamp

TOPICS

[SQL](#)

Related

**CHEAT-SHEET**

SQL Basics Cheat Sheet

**CHEAT-SHEET**

SQL Joins Cheat Sheet

**CHEAT-SHEET**

MySQL Basics Cheat Sheet

[See More →](#)

Grow your data skills with DataCamp for Mobile

Make progress on the go with our mobile courses and daily 5-minute coding challenges.



LEARN

[Learn Python](#)

[Learn R](#)

[Learn AI](#)

[Learn SQL](#)

[Learn Power BI](#)

[Learn Tableau](#)

[Learn Data Engineering](#)

[Assessments](#)

[Career Tracks](#)

[Skill Tracks](#)

[Courses](#)

[Data Science Roadmap](#)

DATA COURSES

[Python Courses](#)

[R Courses](#)

[SQL Courses](#)

[Power BI Courses](#)

[Tableau Courses](#)

[Alteryx Courses](#)

[Azure Courses](#)[Google Sheets Courses](#)[AI Courses](#)[Data Analysis Courses](#)[Data Visualization Courses](#)[Machine Learning Courses](#)[Data Engineering Courses](#)[Probability & Statistics Courses](#)

DATALAB

[Get Started](#)[Pricing](#)[Security](#)[Documentation](#)

CERTIFICATION

[Certifications](#)[Data Scientist](#)[Data Analyst](#)[Data Engineer](#)[SQL Associate](#)[Power BI Data Analyst](#)[Tableau Certified Data Analyst](#)[Azure Fundamentals](#)[AI Fundamentals](#)

RESOURCES

[Resource Center](#)

[Upcoming Events](#)

[Blog](#)

[Code-Alongs](#)

[Tutorials](#)

[Docs](#)

[Open Source](#)

[RDocumentation](#)

[Course Editor](#)

[Book a Demo with DataCamp for Business](#)

[Data Portfolio](#)

[Portfolio Leaderboard](#)

PLANS

[Pricing](#)

[For Business](#)

[For Universities](#)

[Discounts, Promos & Sales](#)

[DataCamp Donates](#)

FOR BUSINESS

[Business Pricing](#)

[Teams Plan](#)

[Data & AI Unlimited Plan](#)

[Customer Stories](#)[Partner Program](#)

ABOUT

[About Us](#)[Learner Stories](#)[Careers](#)[Become an Instructor](#)[Press](#)[Leadership](#)[Contact Us](#)[DataCamp Español](#)[DataCamp Português](#)[DataCamp Deutsch](#)[DataCamp Français](#)

SUPPORT

[Help Center](#)[Become an Affiliate](#)[Privacy Policy](#)[Cookie Notice](#)[Do Not Sell My Personal Information](#)[Accessibility](#)[Security](#)[Terms of Use](#)

